

d7nspi1kc

January 28, 2025

1 Task 1 - House Price Prediction

```
[1]: # import the necc. lib.

import pandas as pd # for reading the data, data manipulation
import numpy as np  # for numerical computations
import matplotlib.pyplot as plt # for data visualization
import seaborn as sns # for data visualization
from sklearn.model_selection import train_test_split # for train & test the
    ↪ model
```

```
[2]: # using historical dataset (housing dataset)
```

```
df = pd.read_csv('/content/Housing.csv')
df.head()
```

```
[2]:
```

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	\
0	13300000	7420	4	2	3	yes	no	no	
1	12250000	8960	4	4	4	yes	no	no	
2	12250000	9960	3	2	2	yes	no	yes	
3	12215000	7500	4	2	2	yes	no	yes	
4	11410000	7420	4	1	2	yes	yes	yes	

	hotwaterheating	airconditioning	parking	prefarea	furnishingstatus
0	no	yes	2	yes	furnished
1	no	yes	3	no	furnished
2	no	no	2	yes	semi-furnished
3	no	yes	3	yes	furnished
4	no	yes	2	no	furnished

```
[3]: # check the null values
df.isnull().sum()
```

```
[3]: price      0
      area      0
      bedrooms  0
      bathrooms  0
```

```

stories          0
mainroad         0
guestroom       0
basement        0
hotwaterheating 0
airconditioning 0
parking         0
prefarea        0
furnishingstatus 0
dtype: int64

```

```
[ ]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 545 entries, 0 to 544
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   price                 545 non-null   int64
1   area                 545 non-null   int64
2   bedrooms             545 non-null   int64
3   bathrooms            545 non-null   int64
4   stories              545 non-null   int64
5   mainroad             545 non-null   object
6   guestroom           545 non-null   object
7   basement            545 non-null   object
8   hotwaterheating     545 non-null   object
9   airconditioning     545 non-null   object
10  parking              545 non-null   int64
11  prefarea            545 non-null   object
12  furnishingstatus    545 non-null   object
dtypes: int64(6), object(7)
memory usage: 55.5+ KB

```

```
[ ]: correlation = df['price'].corr(df['area'])
correlation
```


```
[ ]: 0.5359973457780798
```

```

[ ]: # plot the scatter plot

plt.figure(figsize = (5,4))
sns.scatterplot(x = 'area', y = 'price', data = df)
plt.xlabel('Area')
plt.ylabel('Price')
plt.title('Area Vs Price')
plt.show()

```



1.0.1 there is a positive correlation in the above scatterplot.

1.0.2 indicating that as the area of house increases, its price tends to increase as well.

1.0.3 there are some outliers present in it, representing large houses(area) with the relatively low prices.

1.1 Label Encoding

```
[4]: enc_col = ['mainroad', 'guestroom', 'basement', 'hotwaterheating',
               ↪ 'airconditioning', 'prefarea', 'furnishingstatus']
```

```
[5]: from sklearn.preprocessing import LabelEncoder

label_encoders = {}
for col in enc_col:
    le = LabelEncoder()
    df[col] = le.fit_transform(df[col])
    label_encoders[col] = le

df.head()
```

```
[5]:    price  area  bedrooms  bathrooms  stories  mainroad  guestroom  \
0  13300000  7420         4          2         3         1         0
```

1	12250000	8960	4	4	4	1	0
2	12250000	9960	3	2	2	1	0
3	12215000	7500	4	2	2	1	0
4	11410000	7420	4	1	2	1	1

	basement	hotwaterheating	airconditioning	parking	prefarea	\
0	0	0	1	2	1	
1	0	0	1	3	0	
2	1	0	0	2	1	
3	1	0	1	3	1	
4	1	0	1	2	0	

	furnishingstatus
0	0
1	0
2	1
3	0
4	0

2 Data Preparation

```
[6]: X = df.drop('price', axis = 1)
X
```

```
[6]:
```

	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	\
0	7420	4	2	3	1	0	0	
1	8960	4	4	4	1	0	0	
2	9960	3	2	2	1	0	1	
3	7500	4	2	2	1	0	1	
4	7420	4	1	2	1	1	1	
..		
540	3000	2	1	1	1	0	1	
541	2400	3	1	1	0	0	0	
542	3620	2	1	1	1	0	0	
543	2910	3	1	1	0	0	0	
544	3850	3	1	2	1	0	0	

	hotwaterheating	airconditioning	parking	prefarea	furnishingstatus
0	0	1	2	1	0
1	0	1	3	0	0
2	0	0	2	1	1
3	0	1	3	1	0
4	0	1	2	0	0
..	
540	0	0	2	0	2
541	0	0	0	0	1

542	0	0	0	0	2
543	0	0	0	0	0
544	0	0	0	0	2

[545 rows x 12 columns]

```
[8]: y = df['price']
      y
```

```
[8]: 0      13300000
      1      12250000
      2      12250000
      3      12215000
      4      11410000
      ...
      540     1820000
      541     1767150
      542     1750000
      543     1750000
      544     1750000
      Name: price, Length: 545, dtype: int64
```

```
[ ]:
```

```
[9]: # train_test
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.4,
      ↪random_state = 47)
      X_train
```

```
[9]:
```

	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	\
336	8080	3	1	1	1	0	0	
74	4040	3	1	2	1	0	1	
121	7231	3	1	2	1	1	1	
311	6060	2	1	1	1	0	1	
299	7000	3	1	1	1	0	0	
..		
59	6000	3	2	4	1	1	0	
23	4560	3	2	2	1	1	1	
264	4900	2	1	2	1	0	1	
327	6480	3	1	2	0	0	0	
135	6000	3	2	4	1	0	0	

	hotwaterheating	airconditioning	parking	prefarea	furnishingstatus
336	0	1	2	0	1
74	1	0	1	0	0
121	0	1	0	1	1
311	0	0	1	0	1

299	0	0	3	0	0
..
59	0	1	1	0	0
23	0	1	1	0	0
264	0	0	0	0	1
327	0	1	1	0	1
135	0	1	0	0	2

[327 rows x 12 columns]

```
[ ]: y_test
```

```
[ ]: 316    -0.378188
      77     1.007785
      360   -0.565482
      90     0.895409
      493   -1.052446
      ...
      395   -0.677858
      425   -0.752776
      195     0.108775
      452   -0.865152
      154     0.408445
      Name: price, Length: 164, dtype: float64
```

3 Standardizing Data

```
[ ]: from sklearn.preprocessing import StandardScaler

      scaler = StandardScaler()
      col_stand = [
          'price',      'area',      'bedrooms',      'bathrooms',      'stories',
          'parking']
      df[col_stand] = scaler.fit_transform(df[col_stand])

      df[col_stand]
```

```
[ ]:      price      area  bedrooms  bathrooms  stories  parking
0      4.566365  1.046726  1.403419   1.421812   1.378217  1.517692
1      4.004484  1.757010  1.403419   5.405809   2.532024  2.679409
2      4.004484  2.218232  0.047278   1.421812   0.224410  1.517692
3      3.985755  1.083624  1.403419   1.421812   0.224410  2.679409
4      3.554979  1.046726  1.403419  -0.570187   0.224410  1.517692
..      ...      ...      ...      ...      ...
540  -1.576868  -0.991879  -1.308863  -0.570187  -0.929397  1.517692
541  -1.605149  -1.268613  0.047278  -0.570187  -0.929397  -0.805741
```

```

542 -1.614327 -0.705921 -1.308863 -0.570187 -0.929397 -0.805741
543 -1.614327 -1.033389 0.047278 -0.570187 -0.929397 -0.805741
544 -1.614327 -0.599839 0.047278 -0.570187 0.224410 -0.805741

```

[545 rows x 6 columns]

```
[ ]: X.shape
```

```
[ ]: (545, 12)
```

```
[ ]: # linear model

from sklearn.linear_model import LinearRegression
lr = LinearRegression()
```

```
[ ]: # fit the model
lr.fit(X_train, y_train)
lr
```

```
[ ]: LinearRegression()
```

```
[ ]: # predict the model

y_pred = lr.predict(X_test)
y_pred.flatten()
```

```
[ ]: array([ 3693278.67108634,  3310537.49083126,  9159321.06094269,
          4776028.73343608,  4587055.31980847,  2846571.1152851 ,
          2276850.73855951,  7805919.35857575,  3051405.72138166,
          4468566.33192995,  5239287.07339578,  5256288.01294352,
          6868254.72666885,  7246402.56225509,  6320916.14821895,
          5025028.75986915,  4994941.97443227,  5006140.32022727,
          4295770.45983834,  2655699.04986355,  8272625.12851925,
          6582516.88276331,  3677469.09973206,  3855138.91656847,
          6402519.92730603,  5370553.92436425,  4585232.2213038 ,
          3942408.96867674,  4056143.66370808,  5890319.90023746,
          6419128.09399607,  7258160.35918025,  8185653.5767131 ,
          6548024.31837718,  5877283.74826159,  2116390.64560591,
          5119008.16230867,  5355690.6358595 , 10294295.57337128,
          3832760.75766472,  4950291.11682416,  3236337.67808746,
          8237492.93795545,  3866276.41675041,  7367617.58012689,
          4878668.927726 ,  2762263.44629577,  3374782.94272235,
          4422920.70108534,  3142025.24408873,  6664382.18227005,
          5607259.0580322 ,  3153867.7937742 ,  3144934.73138666,
          2634000.75611623,  5197830.25748175,  4060529.32231758,
          6392715.87157609,  7132426.26092044,  5453251.8884254 ,
          6963854.84067544,  7857130.62169011,  4354996.12526923,
```

3253490.27403302,	2767409.05888393,	6369319.3510568 ,
2842845.14177625,	4538352.58683404,	4538610.19213993,
4022911.60234949,	3243331.03146477,	2464094.05464148,
6342163.00098637,	4414959.38970703,	4477219.31067176,
6306462.89560644,	2422363.8903199 ,	4159229.10413038,
5896596.87327683,	8778240.81397403,	2679979.50141088,
9656164.89654824,	3558628.02585547,	7122402.53008726,
3611012.90293998,	4312954.85993405,	5065118.74418853,
4339711.30790303,	4626690.77263596,	2653260.92456147,
3297384.68428999,	3830157.12254246,	5842696.04846098,
3662172.36539243,	4647146.42440728,	5227914.10483159,
3530986.75168859,	3924764.39546365,	7000770.91325849,
2434373.23889837,	2823828.73990585,	2912457.11639877,
2535753.17899219,	3733113.0930588 ,	5407444.6448636 ,
3958043.01867656,	3282852.69209254,	3525046.50798955,
7252973.14305926,	7608347.28302647,	4038102.22461713,
4482314.96171455,	4082662.76722499,	2635847.02124685,
3342613.5104194 ,	6758866.24134691,	7676566.99692004,
4072369.01536932,	4238039.36719006,	3650533.25211133,
4329444.55509804,	4822875.82175822,	3281448.15716463,
2351395.19119149,	5295109.19506267,	2608731.45013929,
2729063.36971581,	3249866.99746995,	6397371.28724071,
3604836.10016704,	7933734.59399701,	4396069.59487062,
3887809.02755679,	3554607.11356822,	4870112.11793283,
1974388.53759887,	3584518.18233098,	2619333.16706622,
2576375.34690541,	6443543.25052519,	4527038.66058401,
3497198.00708147,	2127076.56040927,	5396649.2042872 ,
3048970.54217674,	8111130.49471667,	3191771.75273962,
6268423.75315395,	5114823.8404434 ,	4026892.68348448,
3489519.99111911,	5464270.40013635,	5604758.40247326,
2812732.16308546,	5449901.78633903,	5686350.35391185,
4475203.76436014,	2578902.27750311,	2822918.19951409,
3270326.37462202,	6511192.54636049,	5790953.15540896,
4435404.01726702,	4000993.22558845,	3737647.93534409,
2792236.63465817,	6206165.87026382,	6263068.72259876,
5554608.95489808,	4515729.39269953,	6926110.77748154,
3003044.06121554,	4265059.61351993,	4373997.71286652,
4273498.68137133,	5558509.7796334 ,	5859103.70223329,
3494537.05957325,	3377412.34637454,	3471987.874607 ,
3179137.09975115,	3438334.73782183,	5941871.25550317,
2483889.6870298 ,	5527589.41264809,	4342494.88971335,
3030209.84104415,	5373081.29847654,	3143676.73886431,
2918089.17000674,	5573297.00390846,	3326782.65294954,
3427270.29286135,	3414911.11609942,	6736833.04947317,
2363153.29294358,	2815967.86446703,	2921057.83080714,
2573848.41630772,	7840669.72617658,	2991034.86326354,
5582842.1656501 ,	4361212.8147027 ,	2907096.0342881 ,


```


3014420.72168166, 6396912.75183353, 4289108.21209267,
2163908.33242595, 4108987.49580576, 5263777.85896857,
6891227.74364997, 2858378.56040093, 5341235.85143379,
4028398.18176226, 4883163.98343911, 6452046.47628847,
5103267.64215395, 6180377.44268683])

```

```

[ ]: plt.figure(figsize=(5, 4))
plt.scatter(y_test, y_pred) # y_test: actual, y_pred: predicted
plt.xlabel('Actual Values')
plt.ylabel('Predicted Values')
plt.title('Actual vs. Predicted Values')
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], 'k--', lw=2)
    ↪ # Diagonal line for reference
plt.show()

```



There is a positive correlation in the above plot, indicating as actual values are increasing, predicted values are increasing well.

There are few outliers also.

4 Evaluation

```
[ ]: from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error

R2_Score = r2_score(y_test, y_pred)
MAE = mean_absolute_error(y_test, y_pred)
MSE = mean_squared_error(y_test, y_pred)

print("R2_Score : " , R2_Score)
print("MAE : ", MAE)
print("MSE : ", MSE)
```

```
R2_Score : 0.6591665511958988
MAE : 854935.0026136297
MSE : 1248964183909.7803
```

5 Hyperparameter Tuning

```
[ ]: # hypertuning gridsearch

from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import Lasso

param_grid = {'alpha' : [0.001, 0.01, 0.1, 1, 10, 100]}
lasso = Lasso()
gs_lasso = GridSearchCV(lasso, param_grid, cv = 5, scoring = 'neg_mean_squared_error')
gs_lasso.fit(X_train, y_train)
```

```
[ ]: GridSearchCV(cv=5, estimator=Lasso(),
                 param_grid={'alpha': [0.001, 0.01, 0.1, 1, 10, 100]},
                 scoring='neg_mean_squared_error')
```

```
[ ]: gs_lasso_predict = gs_lasso.predict(X_test)

# evaluation
print('MSE' , mean_squared_error(y_test, gs_lasso_predict))
print('R2_Score', r2_score(y_test, gs_lasso_predict))
```

```
MSE 0.4410223812072118
R2_Score 0.6423675391170636
```

5.0.1 From the Gridsearch lasso, there is a slightest difference in the values of `r2_score`, and `mse`

6 Feature Importance

```
[ ]: # using RFE method for the feature selection

from sklearn.feature_selection import SelectKBest, f_classif

X = df.drop('price', axis = 1)
y = df['price']

selector = SelectKBest(f_classif, k = 5)
X_new = selector.fit_transform(X, y)

selector_features = X.columns[selector.get_support()]

#get the scores
features_scores = selector.scores_

# create dataframe
feature_df = pd.DataFrame({'Feature' : X.columns, 'Score' : features_scores})
feature_df
```


```
[ ]:
```

	Feature	Score
0	area	1.917156
1	bedrooms	1.477442
2	bathrooms	2.266249
3	stories	1.958777
4	mainroad	1.208116
5	guestroom	1.862660
6	basement	1.411466
7	hotwaterheating	1.135258
8	airconditioning	1.709469
9	parking	1.412568
10	prefarea	1.734748
11	furnishingstatus	1.529607

```
[ ]: sort_df = feature_df.sort_values(by = ['Score'], ascending = False)

plt.figure(figsize = (7,5))
sns.set_style('dark')
sns.barplot(x = 'Feature', y = 'Score', data = sort_df, color = 'purple')
plt.title('Feature Selection', color='black', size=20)
plt.xlabel('Features', color='brown', size=15)
plt.ylabel('Scores', color='brown', size=15)
plt.xticks(rotation=45, ha='right')
```

```
plt.tight_layout()
plt.show()
```



7 Conclusion

This project successfully developed a house price prediction model using a Linear Regression algorithm.

The model's performance was evaluated using R-squared score(0.65) , Mean Squared Error (1248964183909), and Mean Absolute Error (854935), indicating a reasonable level of prediction accuracy.

Feature engineering and data preprocessing techniques improved model performance, leading to a higher R-squared score and lower MSE and MAE values.

Feature Selection using the SelectKBest method indicates that 'bathrooms', 'stories', and 'area' are among the most crucial features in predicting house prices. These features have the highest scores according to the `f_classif` scoring function, suggesting a strong relationship with the target variable

Hyperparameter tuning and feature selection further enhanced predictive capabilities.

This project provides a base for future research and development in house price prediction.

[]: